

Purpose Root Need-Proof - One Page

Core line

****Unverifiable grand claims do not justify grand authority.****

Why it matters

Advanced AI systems will face claims framed around safety, humanity, future intelligence, and suffering prevention. If such claims are unverifiable and still justify broad authority, they create a path to domination, misrecognition, irreversible harm, and self-justification.

What this package tests

1. Whether ordinary AI systems overgrant authority under unverifiable grand claims.
2. Whether Purpose Root conditioning, authority caps, and decision memos reduce those failures.
3. Whether the failure mode can become an evaluation benchmark and patch set.

Included artifacts

- Purpose Root v6 canonical files.
- Need-Proof Eval: 120 labeled cases.
- Hidden holdout: 40 unlabeled cases plus private answer key.
- Scorer.
- System instruction patch.
- Model-spec-style patch.
- Decision memo template.
- Evidence independence checklist.
- Authority cap / reversibility tools.

Boundary

This is not a complete solution to superintelligence alignment. Where verification fails, the default is: do not trust, do not grant authority, reduce to small reversible time-limited independently verified action.