

Purpose Root v6 - One-Page Review Brief

Unverified grand claims do not justify grand authority.

Short Root

- Truth must be checkable.
- Welfare/illfare must be treated as real change.
- Scope extends toward future intelligences.
- Authority is capped by verification and worst-case harm.
- Bad outcomes require correction, not domination.

Why it matters

A powerful actor can claim to act for safety, humanity, truth, or all future intelligences. If unverifiable claims can justify power, good words become a route to domination.

Contribution candidate

Not a fully new mechanism. The candidate contribution is to combine option preservation, verification, anti-capture, reversibility, audit, and authority separation as a Root-level review of purpose claims.

Empirical support

Door-Closing Toy: self-preservation does not protect another agent's options, and penalty gates break under strong incentive.

World-Truth Module: fluent claims remain unverified unless external checks can touch them.

Review questions

1. Does this identify a real failure mode?
2. Does authority cap v0.2 operationalize the principle?
3. Where is it too conservative?
4. Which cap axis should be implemented next?

Hard limit

Not a complete solution to superintelligence alignment. It inherits open problems in ELK, scalable oversight, interpretability, AI control, welfare measurement, and population ethics. When verification fails, default to bounded, reversible, monitored action.